




Pop-up Loft
LONDON

Deep Dive on Amazon RDS (Relational Database Service)


Richard Ainley, AWS Solutions Architect

18th September 2017



Agenda

- Quick intro to RDS
- Security
- Metrics and Monitoring
- High Availability
- Scaling
- Backups and Snapshots
- Migrating



Amazon Relational Database Service (RDS)



Launch

No infrastructure
management




Cost-effective



Application
compatibility



Instant provisioning



Scale up/down



Pop-up Loft
LONDON


Trade-offs with a managed service

Fully managed host and OS

- No access to the database host operating system
- Limited ability to modify configuration that is managed on the host operating system
- No functions that rely on configuration from the host OS

Fully managed storage

- Max storage limits
 - Microsoft SQL Server - 16 TB
 - MySQL, MariaDB, PostgreSQL, Oracle - 6 TB
 - Aurora - 64 TB



Amazon RDS engines

Commercial

ORACLE®



Open source



PostgreSQL



MariaDB

Cloud native



Amazon Aurora





Pop-up Loft
LONDON



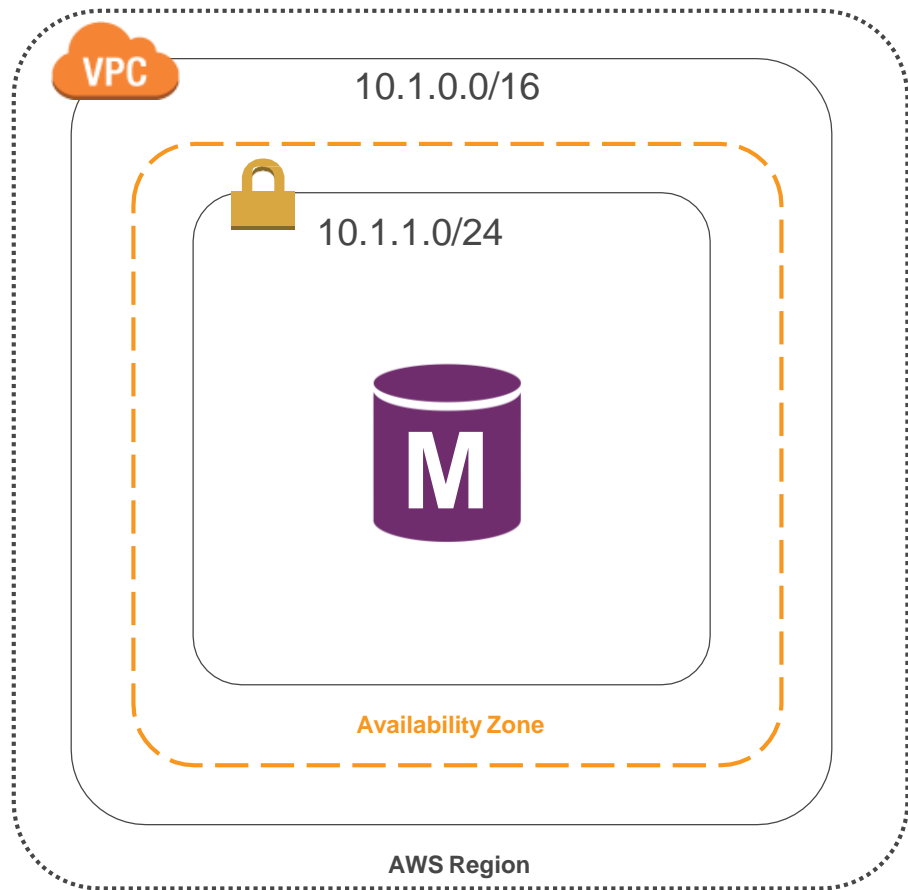
Pop-up Loft
LONDON

Security



Amazon Virtual Private Cloud (Amazon VPC)

Securely control network configuration



Amazon Virtual Private Cloud (Amazon VPC)

Securely control network configuration

Manage connectivity



AWS Direct
Connect



VPN
connection




VPC
peering



Routing
rules

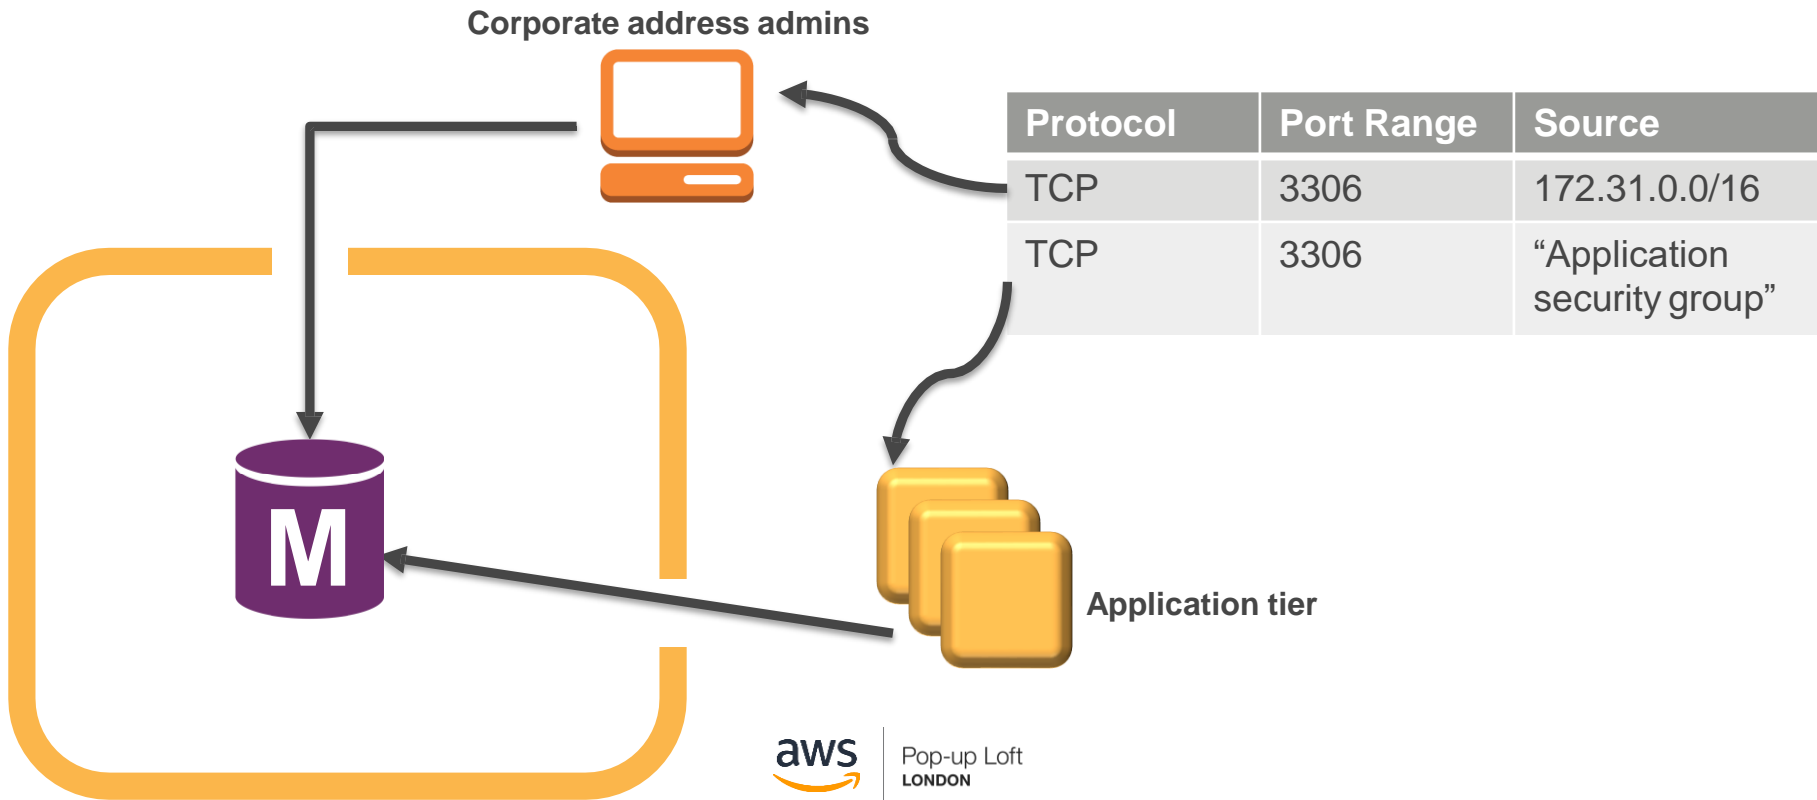


Internet
gateway



Security groups


Database IP firewall protection




AWS IAM governed access

You can use AWS Identity and Access Management (IAM) to:

- Control who can perform actions on RDS
- Authenticate to your RDS MySQL / Aurora DB
 - MySQL 5.6.34 / 5.7.16 or higher
 - Aurora 1.10 or higher
 - Not available for db.t1.micro / db.m1.small



Compliance



Singapore MTCS



27001/9001
27017/27018



Pop-up Loft
LONDON

Compliance

Aurora

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
HIPAA BAA

MySQL

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
FedRamp
HIPAA BAA
UK Gov. Programs
Singapore MTCS

Oracle

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
FedRamp
HIPAA BAA
UK Gov. Programs
Singapore MTCS

PostgreSQL

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
UK Gov. Programs
Singapore MTCS
HIPAA BAA

MariaDB

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
HIPAA BAA


SQL Server

SOC 1, 2, 3
ISO 20001/9001
ISO 27107/27018
PCI
UK Gov. Programs
Singapore MTCS
HIPAA BAA

SSL

Database traffic encryption

Available for all six engines





At Rest Encryption

Transparent Data Encryption



Pop-up Loft
LONDON

At Rest Encryption for all RDS Engines


AWS Key Management Service (KMS)

Two-tiered key hierarchy using envelope encryption:

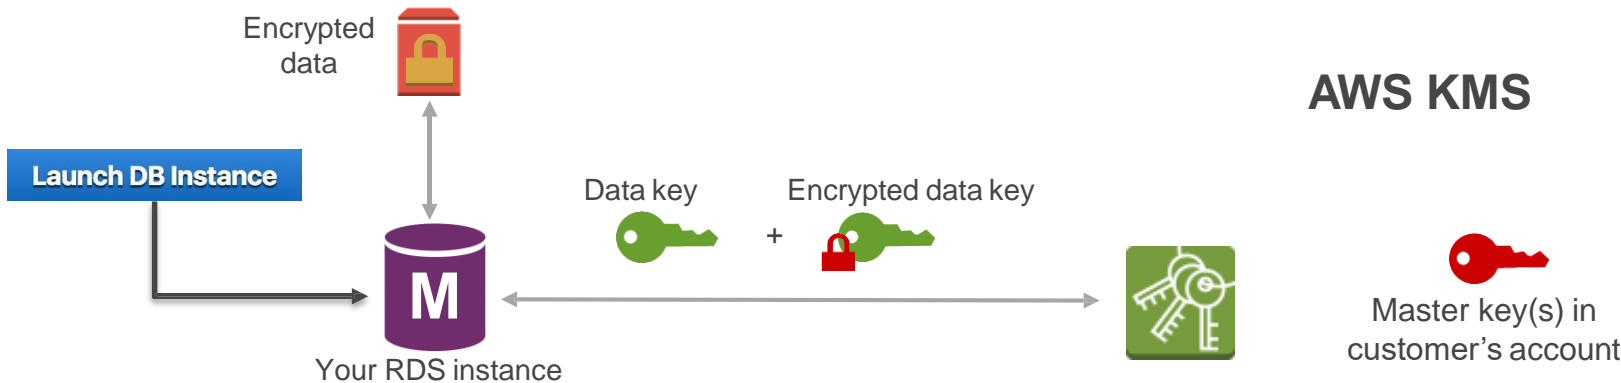
- Unique data key encrypts customer data
- AWS KMS master keys encrypt data keys
- Available for **ALL** RDS engines

Benefits:

- Limits risk of compromised data key
- Better performance for encrypting large data
- Easier to manage small number of master keys than millions of data keys
- Centralized access and audit of key activity



How keys are used to protect your data



1. Launch your RDS instance
2. RDS instance requests encryption key to use to encrypt data, passes reference to master key in account
3. Client request authenticated based on permissions set on both the user and the key
4. A unique data encryption key is created and encrypted under the KMS master key
5. Plaintext and encrypted data key returned to RDS
6. Plaintext data key stored in memory and used to encrypt/decrypt RDS data

Enabling encryption

(default) aws/rds
✓ Enter a key ARN

Console

Enable Encryption	Yes
Master Key	(default) aws/rds
Description	Default master key that protects my RDS database volumes when no other key is defined
Account	This account ()
KMS Key ID	()

AWS Command Line Interface (AWS CLI)



```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test \  
--allocated-storage 20 --storage-encrypted \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```

```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test1 \  
allocated-storage 20 --storage-encrypted --kms-key-id xxxxxxxxxxxxxxxxxxxx \  
instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-  
password myawsuser
```

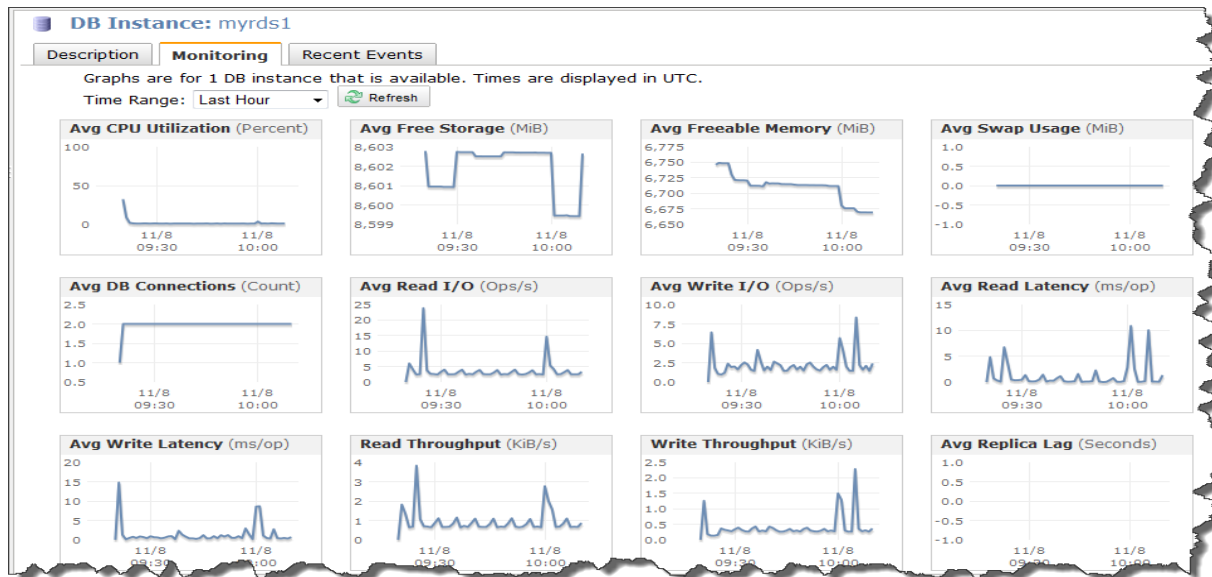


Pop-up Loft
LONDON

Metrics and Monitoring



Standard monitoring



Amazon CloudWatch metrics for Amazon RDS

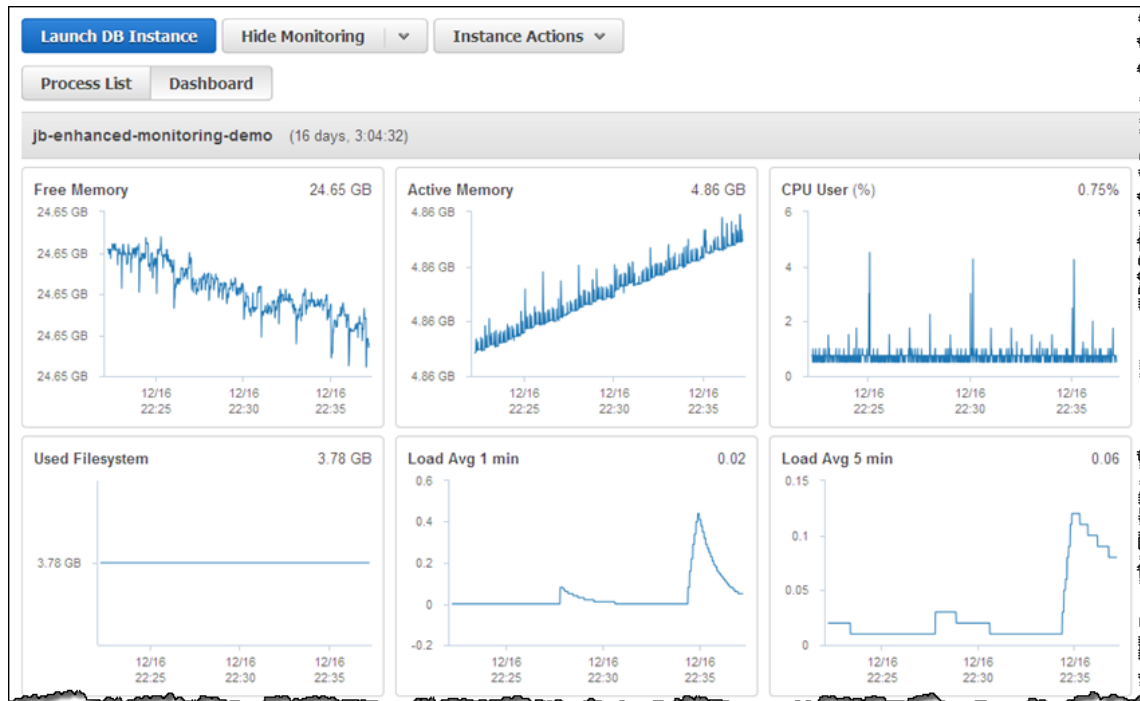
- CPU utilization
- Storage
- Memory
- Swap usage
- DB connections
- I/O (read and write)
- Latency (read and write)
- Throughput (read and write)
- Replica lag
- Many more

Amazon CloudWatch Alarms

- Similar to on-premises custom monitoring tools

Enhanced Monitoring

Access to over 50 new CPU, memory, file system, and disk I/O metrics as low as 1 second intervals



Monitoring

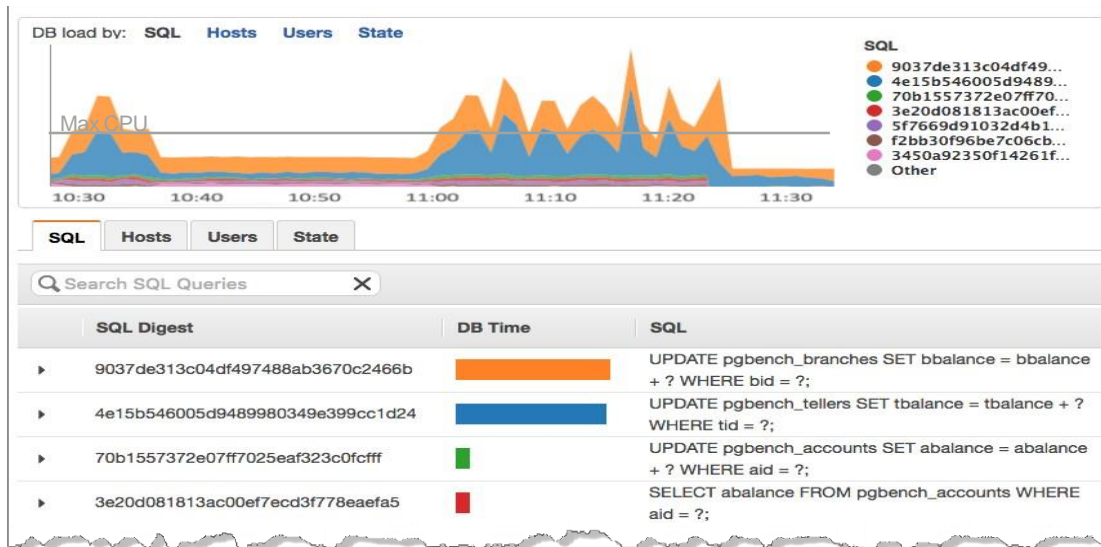
Enable Enhanced Monitoring Yes ▾

Monitoring Role Default ▾

Granularity 1 ▾ second(s)

☒ I authorize RDS to create the IAM role rds-monitoring-role.

Simplify monitoring with AWS Management Console



Amazon Performance Insights for RDS



- Database Load : Identifies database bottlenecks
 - Easy
 - Powerful
- Identifies source of bottlenecks
 - Top SQL
- Adjustable Time frame
 - Hour, day, week and longer

AWS re:Invent 2016 DAT206: https://youtu.be/ztmtJJTC8_Y?t=39m53s

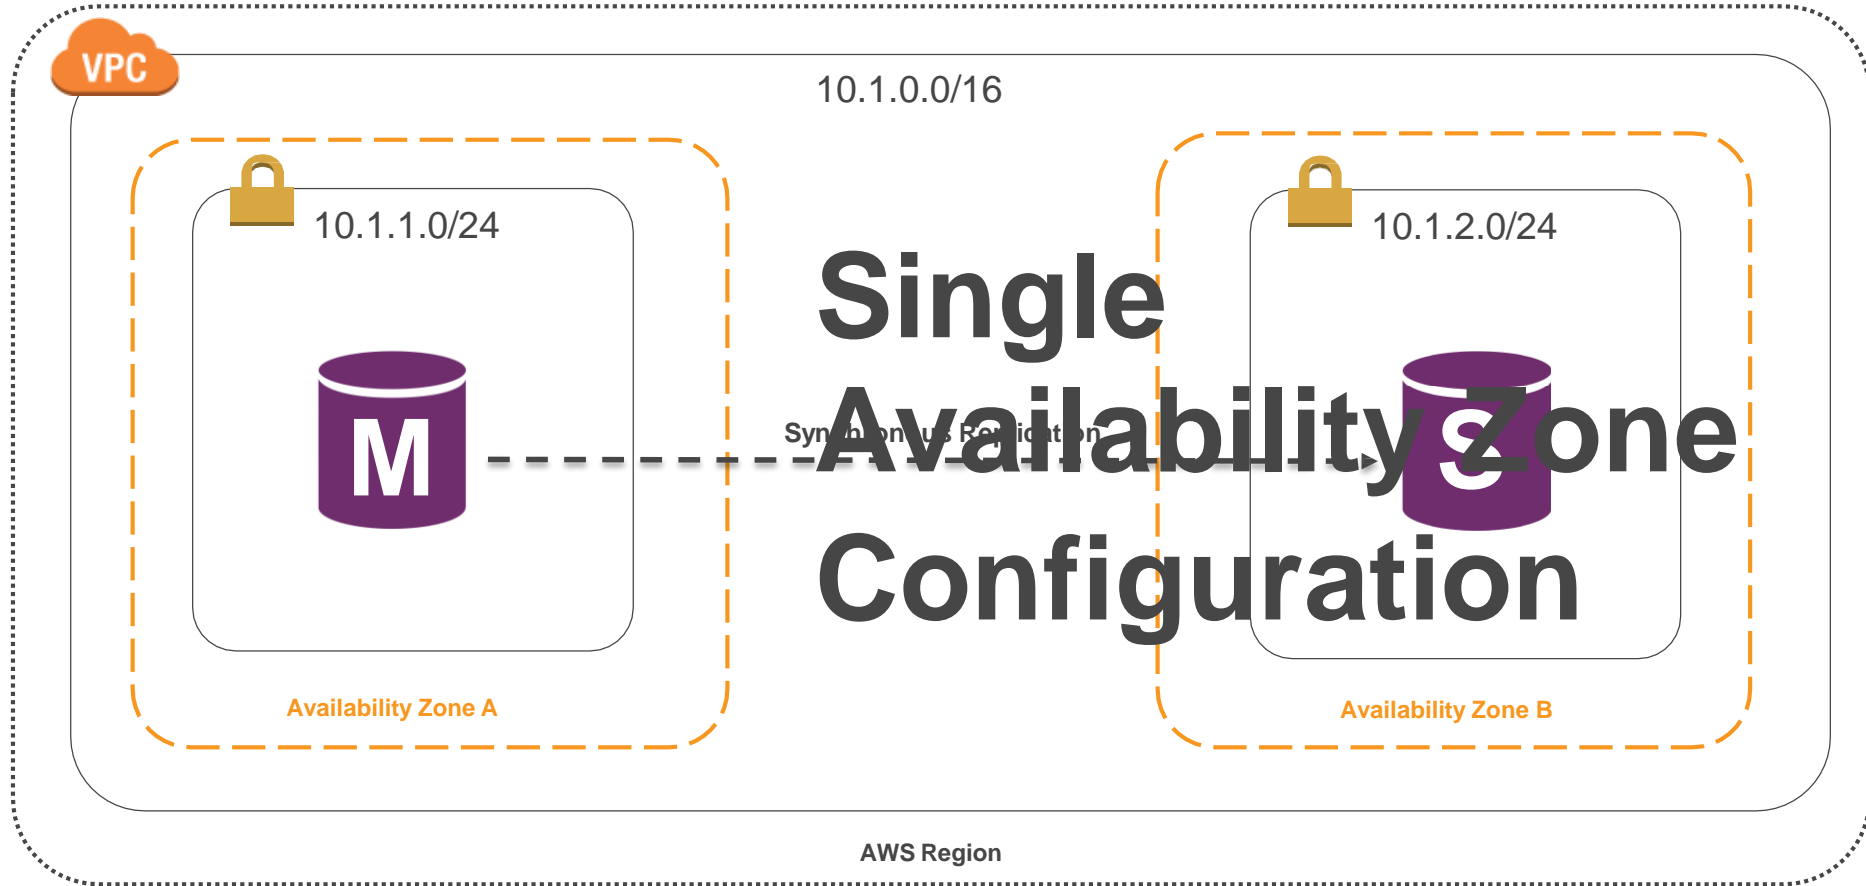


Pop-up Loft
LONDON

High Availability




HA Multi Availability Zone Configuration



High availability - Multi-AZ - DNS

mydatabase.us-west-2.rds.amazonaws.com:3306



Read Replicas

Bring data close to your customer's applications in different regions

Relieve pressure on your master node for supporting reads and writes

Promote a Read Replica to a master for faster recovery in the event of disaster

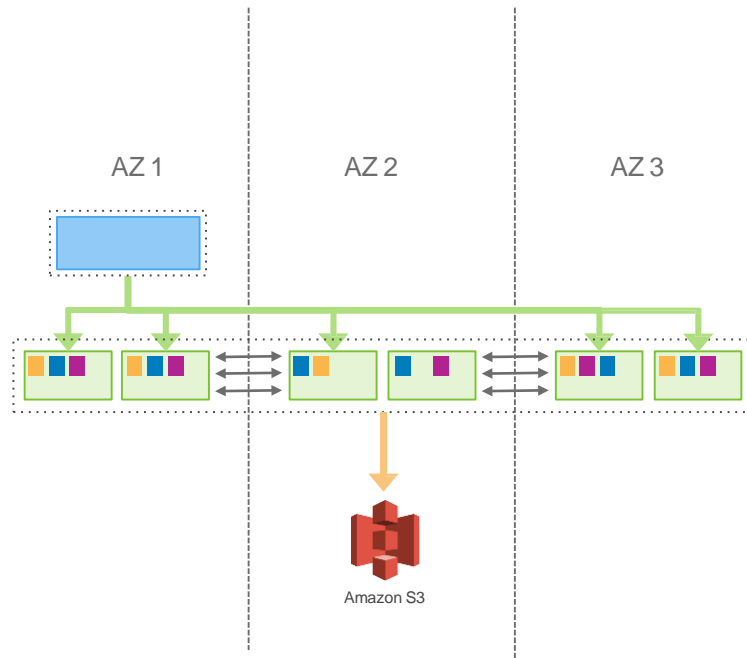


Within or cross-region

- MySQL
- MariaDB
- PostgreSQL
- Aurora

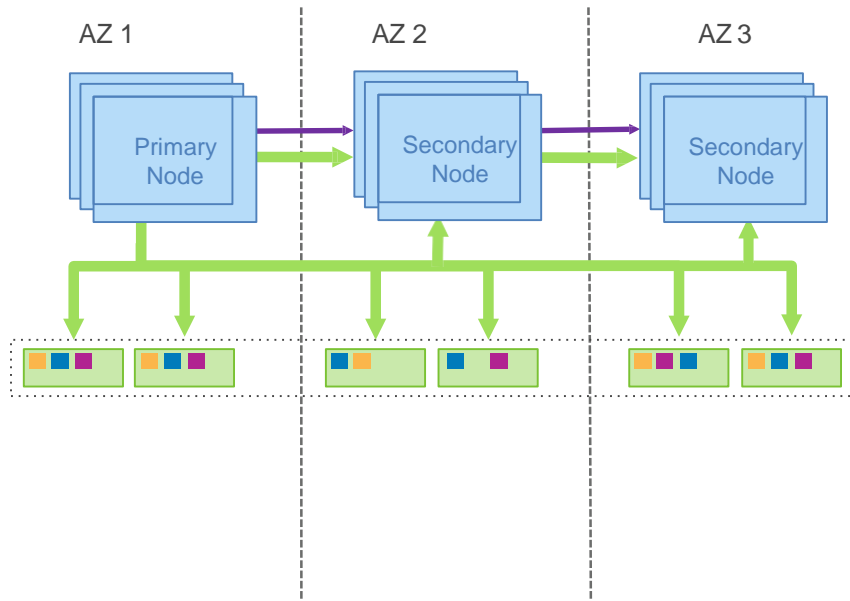
High availability - Amazon Aurora storage

- Storage volume automatically grows up to 64 TB
- Quorum system for read/write; latency tolerant
- Peer-to-peer gossip replication to fill in holes
- Continuous backup to Amazon S3 (built for 11 9s durability)
- Continuous monitoring of nodes and disks for repair
- 10 GB segments as unit of repair or hotspot rebalance
- Quorum membership changes do not stall writes



High availability - Amazon Aurora nodes



- Aurora cluster contains primary node and up to 15 secondary nodes
- Failing database nodes are automatically detected and replaced
- Failing database processes are automatically detected and recycled
- Secondary nodes automatically promoted on persistent outage, no single point of failure
- Customer application can scale out read traffic across secondary nodes






Pop-up Loft
LONDON

Scaling



Why Scale?

- Handle higher load or lower usage
- Naturally grow over time
- Control costs



What can I scale?

Database Instance



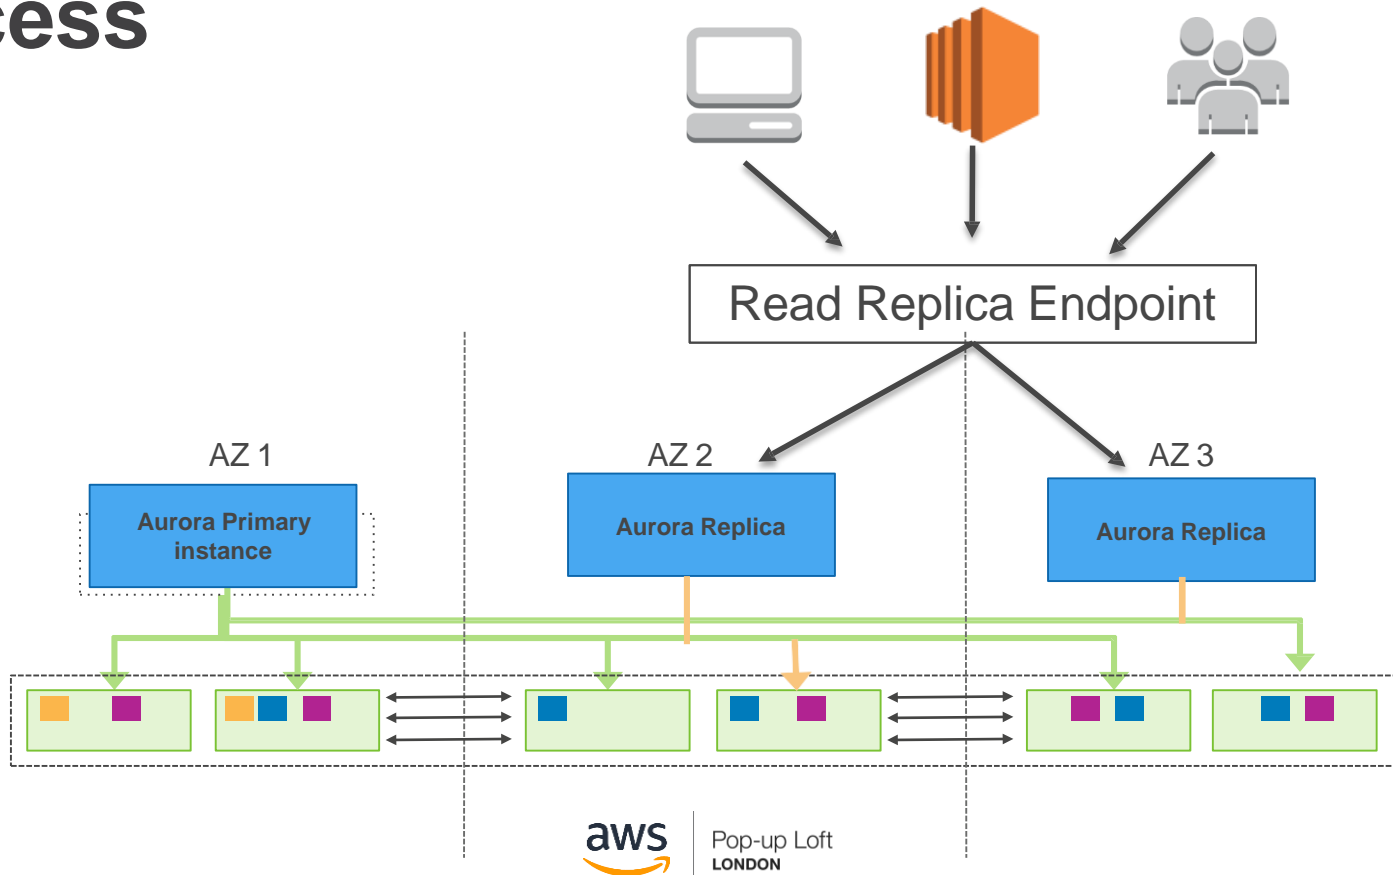
Read Replicas



Storage



Amazon Aurora - Balanced Read Replica Access



Scaling your instance up/down

AWS Management Console

Instance Actions ▾

- See Details
- Create Read Replica
- Promote Read Replica
- Take Snapshot
- Restore to Point in Time
- Migrate Latest Snapshot
- Modify**
- Reboot
- Delete

Modify DB Instance: mysql-test

Instance Specifications

DB Engine Version	MySQL 5.6.27 (default) ▾
DB Instance Class	db.m4.xlarge — 4 vCPU, 16 GiB RAM ▾
Multi-AZ Deployment	Yes ▾
Storage Type	General Purpose (SSD) ▾
Allocated Storage*	1600 GB

 **Apply Immediately**  

Scaling - single Availability Zone

mydatabase.us-west-2.rds.amazonaws.com:3306

Alarms and Recent Events

TIME (UTC-7)	EVENT
Oct 30 8:06 AM	Finished applying modification to DB instance class
Oct 30 8:06 AM	DB instance restarted
Oct 30 8:03 AM	DB instance shutdown
Oct 30 8:00 AM	Applying modification to database instance class

AWS Region

Scaling - Multi-AZ

Alarms and Recent Events

TIME (UTC-7)	EVENT
Oct 30 8:11 AM	Finished applying modification to DB instance class
Oct 30 8:03 AM	Multi-AZ instance failover completed
Oct 30 8:02 AM	DB instance restarted
Oct 30 8:02 AM	Multi-AZ instance failover started
Oct 30 7:52 AM	Applying modification to database instance class

onaws.com:3306



AWS Region

Scaling - automation

AWS CLI

```
aws rds modify-db-instance --db-instance-identifier sg-cli-test --db-instance-class db.m4.large --apply-immediately
```

Scheduled CLI—cron

```
#Scale down at 8:00 PM on Friday  
0 20 * * 5 /home/ec2-user/scripts/scale_down_rds.sh  
  
#Scale up at 4:00 AM on Monday  
0 4 * * 1 /home/ec2-user/scripts/scale_up_rds.sh
```

Scaling - automation

Scheduled - AWS Lambda

No server but still runs on a schedule!

```
import boto3

client=boto3.client('rds')


def lambda_handler(event, context):
    response=client.modify_db_instance(DBInstanceIdentifier='sg-cli-test',
                                       DBInstanceClass='db.m4.xlarge',
                                       ApplyImmediately=True)

    print response
```

Scaling - automation

Metrics-based scaling

- Amazon CloudWatch and AWS Lambda!



Scaling - automation

```
import boto3
import json

client=boto3.client('rds')


def lambda_handler(event, context):

    message = event['Records'][0]['Sns']['Message']
    parsed_message=json.loads(message)
    db_instance=parsed_message['Trigger']['Dimensions'][0]['value']
    print 'DB Instance: ' + db_instance

    response=client.modify_db_instance(DBInstanceIdentifier=db_instance,
                                       DBInstanceClass='db.m4.large',
                                       ApplyImmediately=True)

    print response
```

Switch Off Dev Test Instances





MySQL
MariaDB
PostgreSQL
Oracle
SQL Server



Pop-up Loft
LONDON

Backups and snapshots



Automated Backups

MySQL, PostgreSQL, MariaDB, Oracle, SQL Server

- Scheduled daily volume backup of entire instance
- Archive database change logs
- 35-day retention
- Taken from standby when running multi-AZ

Aurora

- Automatic, continuous, incremental backups
- No impact on database performance
- 35-day retention



How do automated RDS backups work?

Every day during your backup window, the RDS service creates a storage volume snapshot of your database

→ If database is Multi-AZ, the snapshot is taken from the standby

Every five minutes, RDS backs up the transaction logs of your database

Availability and Durability

DB Instance Status available

Multi AZ Yes

Automated Backups Enabled (7 Days)

Latest Restore Time October 12, 2016 at 4:50:00 PM UTC-7



Pop-up Loft
LONDON

Restoring

- Creates an entire new database instance
- You define all the instance configuration, just like creating a new instance

Restore DB Instance

You are creating a new DB Instance from a source DB Instance at a specified time. This new DB Instance will have the default DB Security Group and DB Parameter Groups.

Use Latest Restorable Time ☒ October 30, 2016 at 7:49:24 AM UTC-7


Use Custom Restore Time ☐ MMMM d, y : : UTC-7



Pop-up Loft
LONDON

Snapshots



- Full copies of your RDS database
- Independent of scheduled backups
- Used to create a new RDS instance
- Taken from the standby when running multi-AZ



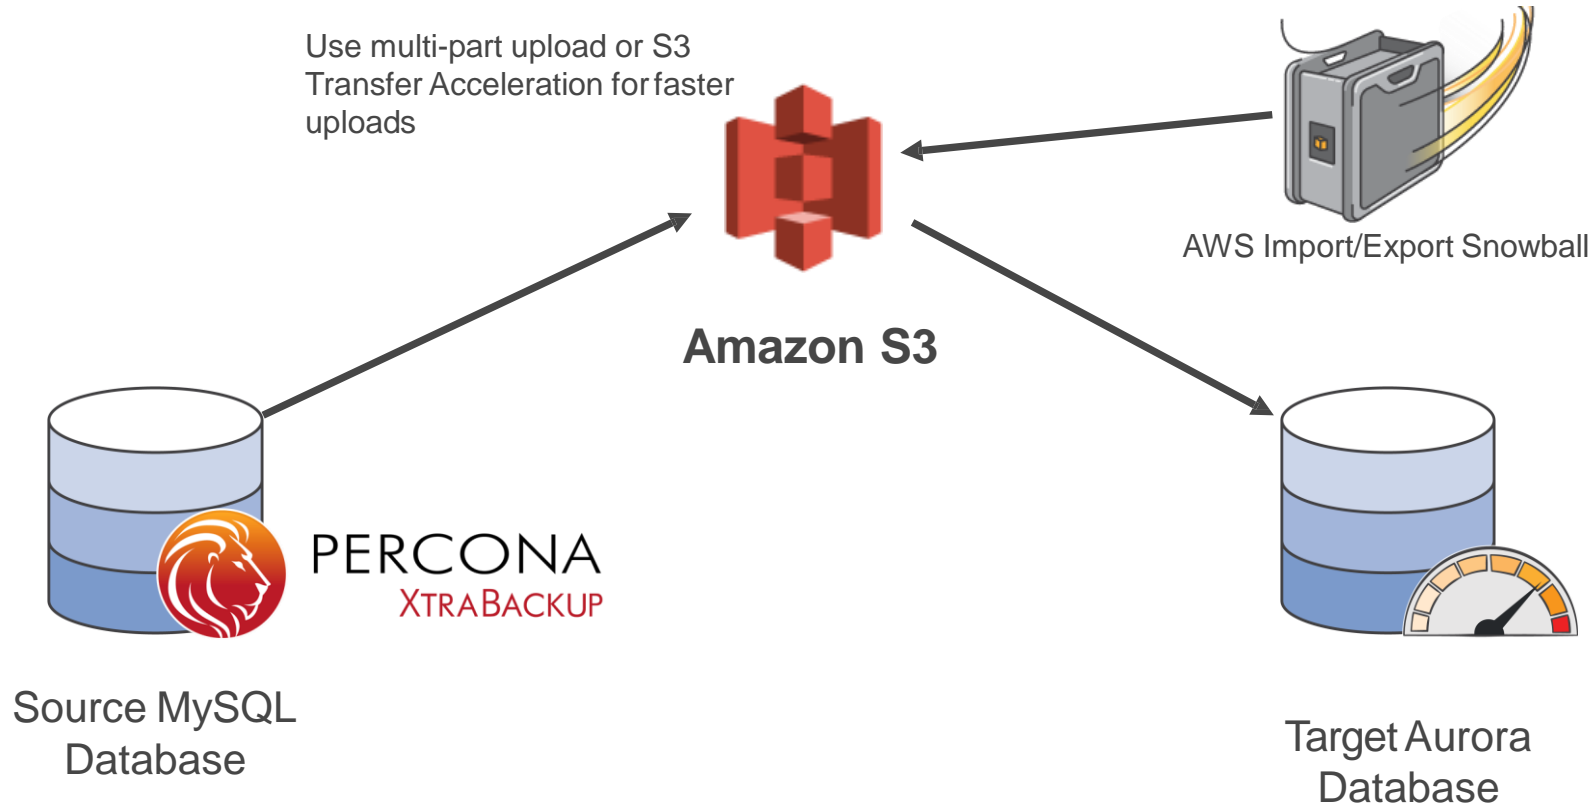


Pop-up Loft
LONDON

Migrating onto RDS



MySQL Backup to Aurora via S3



MySQL Backup to Aurora via S3

Restore Aurora DB Cluster from S3



Specify Source Backup Details

Source Database Specifications

Source Engine

mysql

Source Engine Version

5.6

S3 Backup Location

Select S3 Bucket*

mysql-demo-backups

S3 Bucket Prefix (Optional)

IAM Role

IAM Role*

db-backup-file-access

Create a New Role

* Required


Cancel

Next Step



Pop-up Loft
LONDON

SQL Server Backup to RDS SQL Server via S3



SQL Server Backup to RDS SQL Server via S3

Importing to RDS

No Encryption

```
exec msdb.dbo.rds_restore_database  
    @restore_db_name='database_name',  
    @s3_arn_to_restore_from='arn:aws:s3:::bucket_name/file_name_and_extension';
```

Encryption

```
exec msdb.dbo.rds_restore_database  
    @restore_db_name='database_name',  
    @s3_arn_to_restore_from='arn:aws:s3:::bucket_name/file_name_and_extension',  
    @kms_master_key_arn='arn:aws:kms:region:account-id:key/key-id';
```



AWS Database Migration Service




ORACLE

Amazon Aurora



- ✓ Move data to the same or different database engine
- ✓ Keep your apps running during the migration
- ✓ Start your first migration in 10 minutes or less
- ✓ Replicate within, to, or from Amazon EC2 or RDS



AWS Schema Conversion Tool

- ✓ Migrate from Oracle and SQL Server
- ✓ Move your tables, views, stored procedures, and data manipulation language (DML) to MySQL, MariaDB, and Aurora
- ✓ Highlight where manual edits are needed

SCT supported OLTP conversions

Source Database	Target Database on Amazon RDS
Microsoft SQL Server (version 2008 and later)	Amazon Aurora (MySQL or PostgreSQL), Microsoft SQL Server, MySQL, PostgreSQL
MySQL (version 5.5 and later)	Amazon Aurora (PostgreSQL), MySQL, PostgreSQL
Oracle (version 10.2 and later)	Amazon Aurora (MySQL or PostgreSQL), MySQL, Oracle, PostgreSQL
PostgreSQL (version 9.1 and later)	Amazon Aurora (MySQL), MySQL, PostgreSQL



Pop-up Loft
LONDON

Thank you!

