

# Apache Parquet


# Apache Parquet

- *Apache Parquet is an binary*
- *Efficient columnar data format*
- uses various techniques to store data in a CPU and I/O efficient way
- row groups, compression for pages in column chunks or dictionary encoding for columns
- Index hints and statistics to quickly skip over chunks of irrelevant data enable efficient queries on large amount of data
- *Supported in Hadoop, Spark, Pandas, AWS, Azure Big Data Platforms*


	day	location	product	sale
row 1	2017-01-01	l1	p1	300
row 2	2017-01-01	l1	p2	40
row 3	2017-01-01	l2	p1	44
row 4	2017-02-01	l1	p1	200

Traditional Memory Buffer	
row 1	2017-01-01
	l1
	p1
	300
row 2	2017-01-01
	l1
	p2
	40
row 3	2017-01-01
	l2
	p1
	44

Columnar Storage	
day	2017-01-01
	2017-01-01
	2017-01-01
	2017-01-02
location	l1
	l1
	l2
	l1
product	p1
	p2
	p1
	p1



# Columnar storage




Nested schema

Logical table representation



a	b	c
a1	b1	c1
a2	b2	c2
a3	b3	c3
a4	b4	c4
a5	b5	c5

On Disk:

Row layout



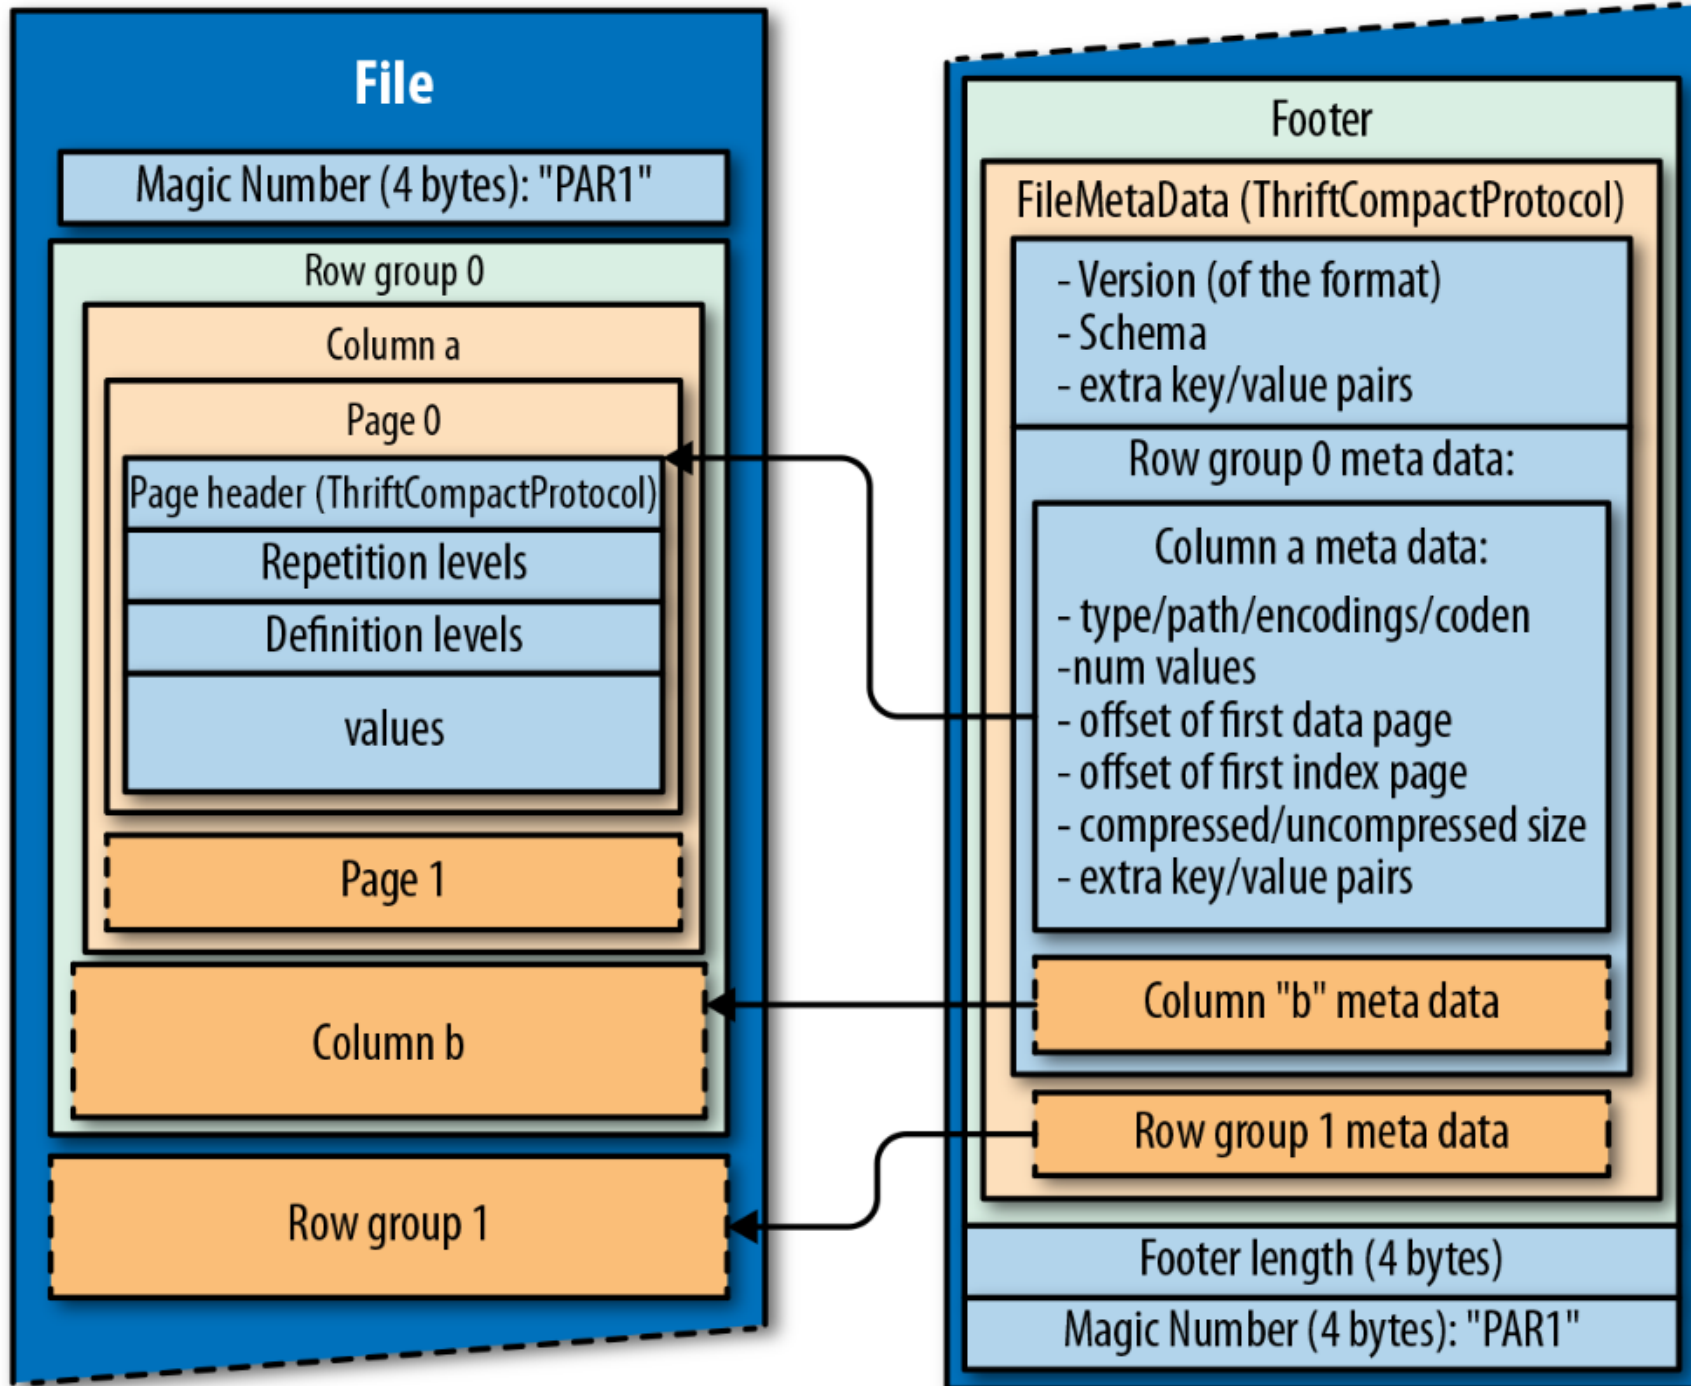
Column layout



Encodings: Dictionary, RLE, Delta, Prefix

## Key Concepts

Block Size  
Row Group  
Page



# Unit of parallelization

- MapReduce - File/Row Group
- IO - Column chunk
- Encoding/Compression - Page

# Compression

- Supports (page) compression and splitting
- Supports nested columns (Dremel encoding)
- Supports minimal number of types

# Glossary

- Block (hdfs block): This means a block in hdfs and the meaning is unchanged for describing this file format. The file format is designed to work well on top of hdfs.
- File: A hdfs file that must include the metadata for the file. It does not need to actually contain the data.
- Row group: A logical horizontal partitioning of the data into rows. There is no physical structure that is guaranteed for a row group. A row group consists of a column chunk for each column in the dataset.
- Column chunk: A chunk of the data for a particular column. These live in a particular row group and is guaranteed to be contiguous in the file.
- Page: Column chunks are divided up into pages. A page is conceptually an indivisible unit (in terms of compression and encoding). There can be multiple page types which is interleaved in a column chunk.
- Hierarchically, a file consists of one or more row groups. A row group contains exactly one column chunk per column. Column chunks contain one or more pages



# File Format

- 4-byte magic number "PAR1"
- <Column 1 Chunk 1 + Column Metadata>
- <Column 2 Chunk 1 + Column Metadata>
- ...
- <Column N Chunk 1 + Column Metadata>
- <Column 1 Chunk 2 + Column Metadata>
- <Column 2 Chunk 2 + Column Metadata>
- ...
- <Column N Chunk 2 + Column Metadata>
- ...
- <Column 1 Chunk M + Column Metadata>
- <Column 2 Chunk M + Column Metadata>
- ...
- <Column N Chunk M + Column Metadata>
- File Metadata
- 4-byte length in bytes of file metadata
- 4-byte magic number "PAR1"

There are N columns in this table, split into M row groups. The file metadata contains the locations of all the column metadata start locations. More details on what is contained in the metadata can be found in the thrift files.

Metadata is written after the data to allow for single pass writing.

Readers are expected to first read the file metadata to find all the column chunks they are interested in

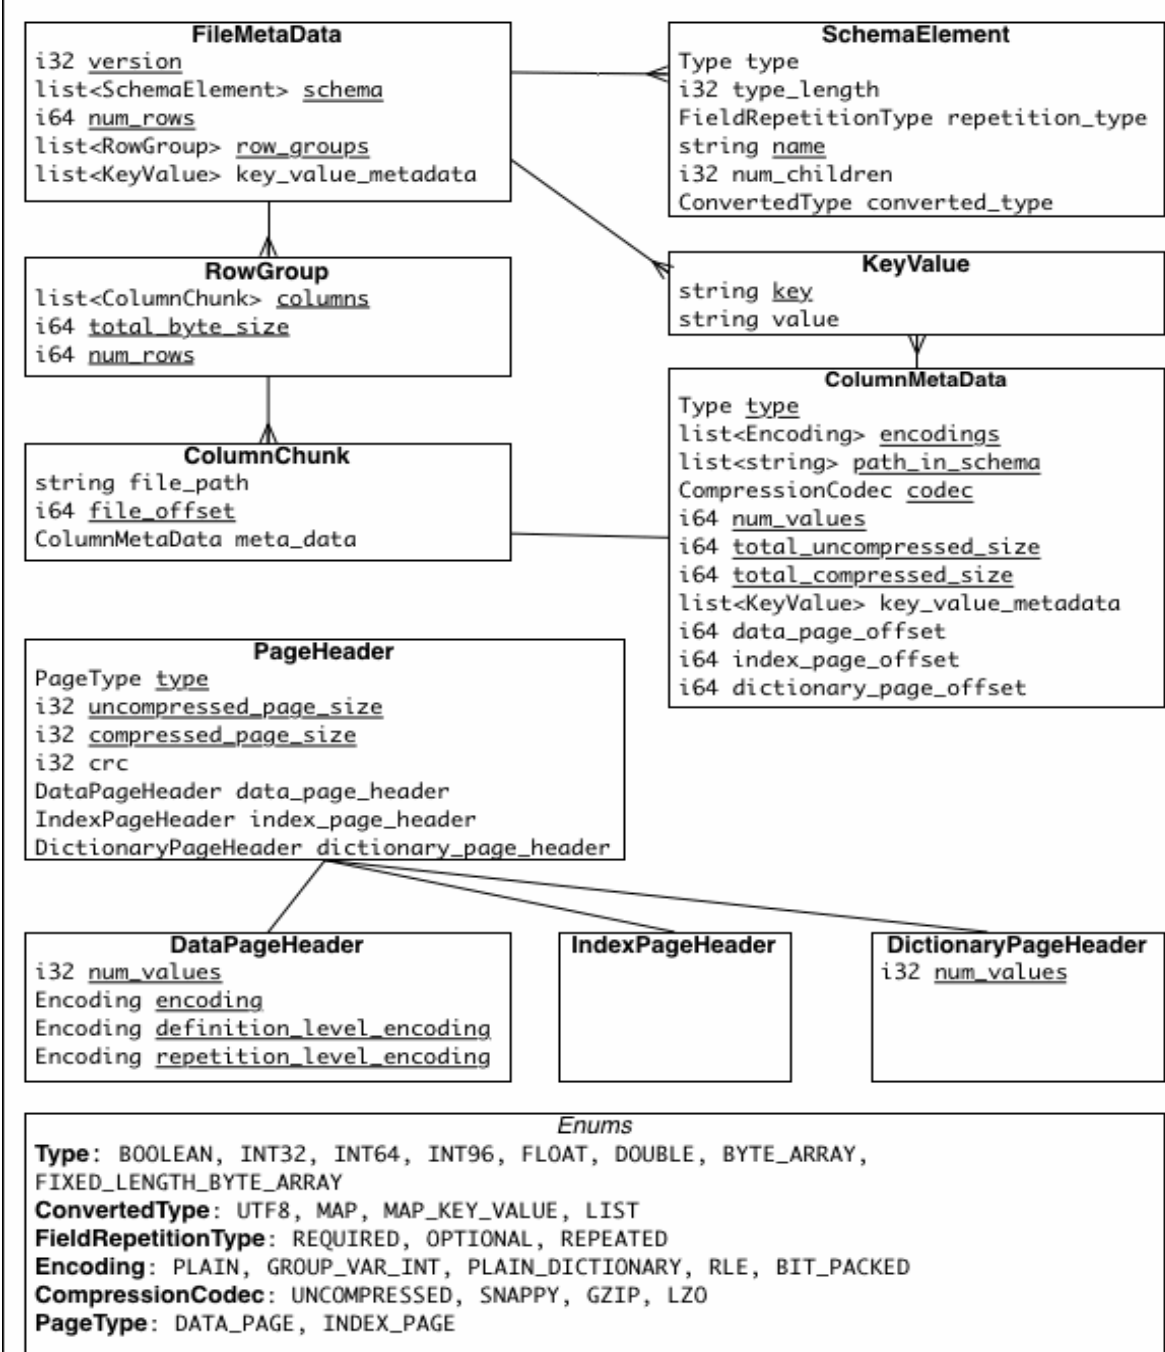
The columns chunks should then be read sequentially

# Meta Data

## three types of metadata

- file metadata
- column (chunk) metadata
- page header metadata

All thrift structures are serialized using the TCompactProtocol



# Types

- BOOLEAN: 1 bit boolean
- INT32: 32 bit signed ints
- INT64: 64 bit signed ints
- INT96: 96 bit signed ints
- FLOAT: IEEE 32-bit floating point values
- DOUBLE: IEEE 64-bit floating point values
- BYTE\_ARRAY: arbitrarily long byte arrays.

16-bit ints are not explicitly supported in the storage format since they are covered by 32-bit ints with an efficient encoding

# Logical Types

- Logical types are used to extend the types that parquet can be used to store
- strings are stored as byte arrays (binary) with a UTF8 annotation
- Annotations are stored as a ConvertedType in the file metadata
- STRING, ENUM, UUID, SIGNED INT, UNSIGNED INT, DECIMAL, DATE, TIME, TIMESTAMP
- For more <https://github.com/apache/parquet-format/blob/master/LogicalTypes.md>

# Reference

- <http://parquet.apache.org/documentation/latest/>