

facebook

HIVE

Data Warehousing & Analytics on Hadoop

Facebook Data Team


Why Another Data Warehousing System?

- Problem: Data, data and more data
 - 200GB per day in March 2008
 - 2+TB(compressed) raw data per day today
- The Hadoop Experiment
 - Much superior to availability and scalability of commercial DBs
 - Efficiency not that great, but throw more hardware
 - Partial Availability/resilience/scale more important than ACID
- Problem: Programmability and Metadata
 - Map-reduce hard to program (users know sql/bash/python)
 - Need to publish data in well known schemas
- Solution: HIVE

What is HIVE?

- A system for querying and managing structured data built on top of Hadoop
 - Uses Map-Reduce for execution
 - HDFS for storage - but any system that implements Hadoop FS API
- Key Building Principles:
 - Structured data with rich data types (structs, lists and maps)
 - Directly query data from different formats (text/binary) and file formats (Flat/Sequence)
 - SQL as a familiar programming tool and for standard analytics
 - Allow embedded scripts for extensibility and for non standard applications
 - Rich MetaData to allow data discovery and for optimization

Data Warehousing at Facebook Today



Hive/Hadoop Usage @ Facebook

- Types of Applications:
 - Summarization
 - Eg: Daily/Weekly aggregations of impression/click counts
 - Complex measures of user engagement
 - Ad hoc Analysis
 - Eg: how many group admins broken down by state/country
 - Data Mining (Assembling training data)
 - Eg: User Engagement as a function of user attributes
 - Spam Detection
 - Anomalous patterns in UGC
 - Application api usage patterns
 - Ad Optimization
 - Too many to count ..

Hadoop Usage @ Facebook

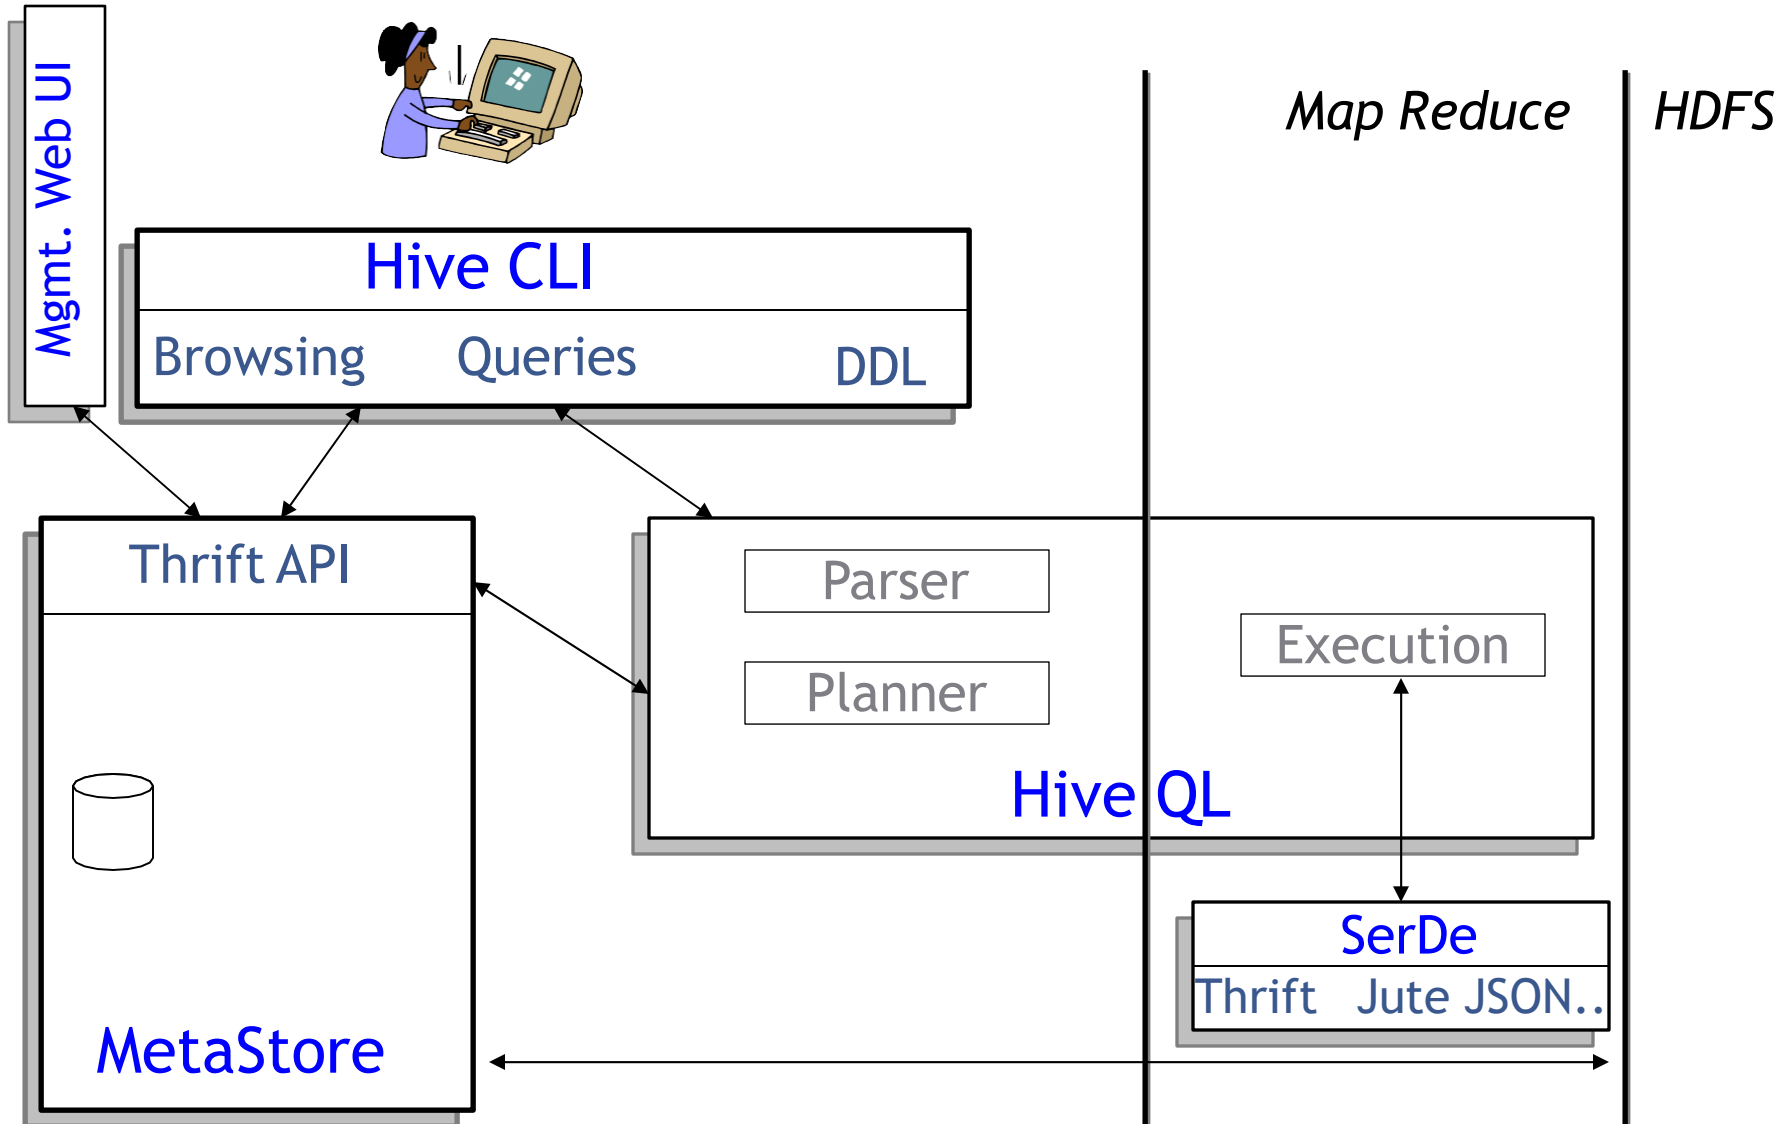
- Data statistics:

- Total Data: 180TB (mostly compressed)
- Net Data added/day: 2+TB (compressed)
 - 6TB of uncompressed source logs
 - 4TB of uncompressed dimension data reloaded daily

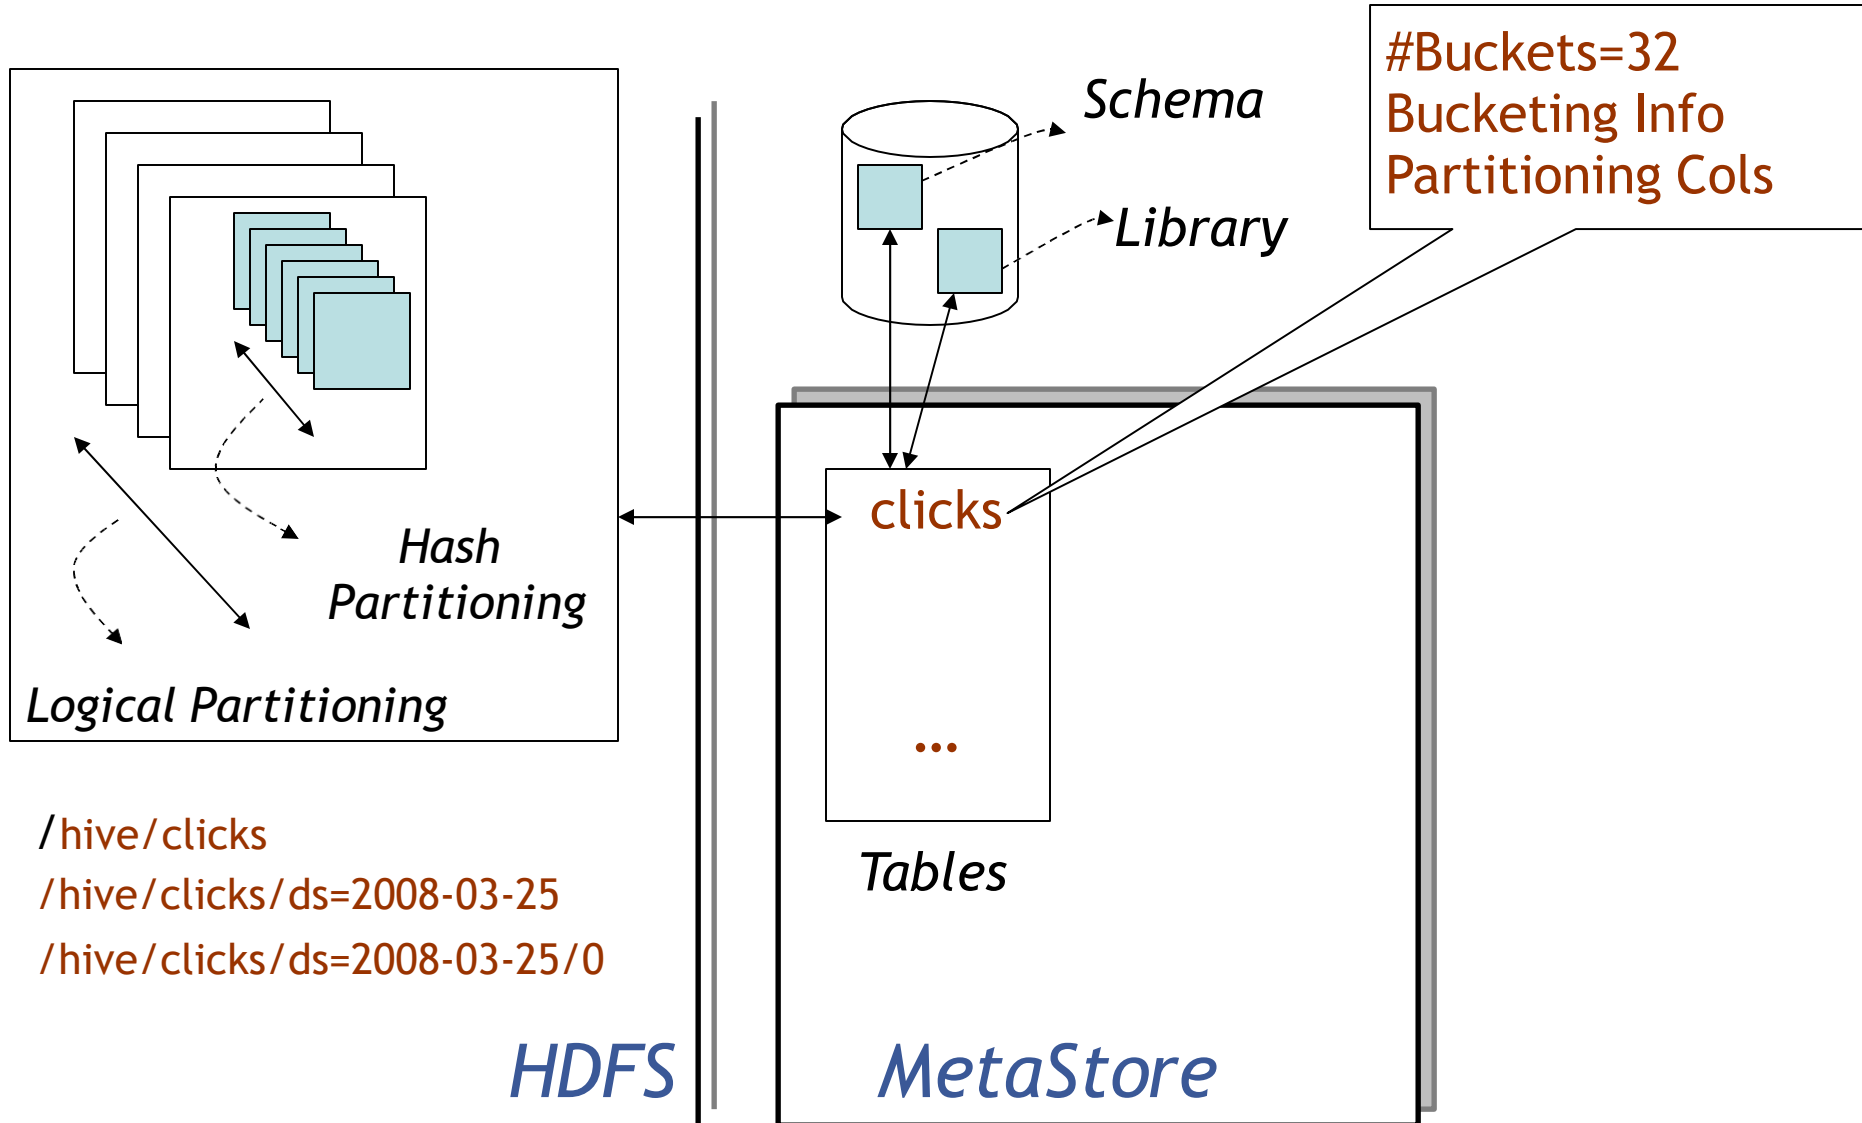
- Usage statistics:

- 3200 jobs/day with 800K tasks(map-reduce tasks)/day
- 55TB of compressed data scanned daily
- 15TB of compressed output data written to hdfs
- 80 MM compute minutes/day

HIVE: Components



Data Model



Dealing with Structured Data

- Type system
 - Primitive types
 - Recursively build up using Composition/Maps/Lists
- ObjectInspector interface for user-defined types
 - To recursively list schema
 - To recursively access fields within a row object
- Generic (De)Serialization Interface (SerDe)
- Serialization families implement interface
 - Thrift DDL based SerDe
 - Delimited text based SerDe
 - You can write your own SerDe (XML, JSON ...)

MetaStore

- Stores Table/Partition properties:
 - Table schema and SerDe library
 - Table Location on HDFS
 - Logical Partitioning keys and types
 - Partition level metadata
 - Other information
- Thrift API
 - Current clients in Php (Web Interface), Python interface to Hive, Java (Query Engine and CLI)
- Metadata stored in any SQL backend
- Future
 - Statistics
 - Schema Evolution

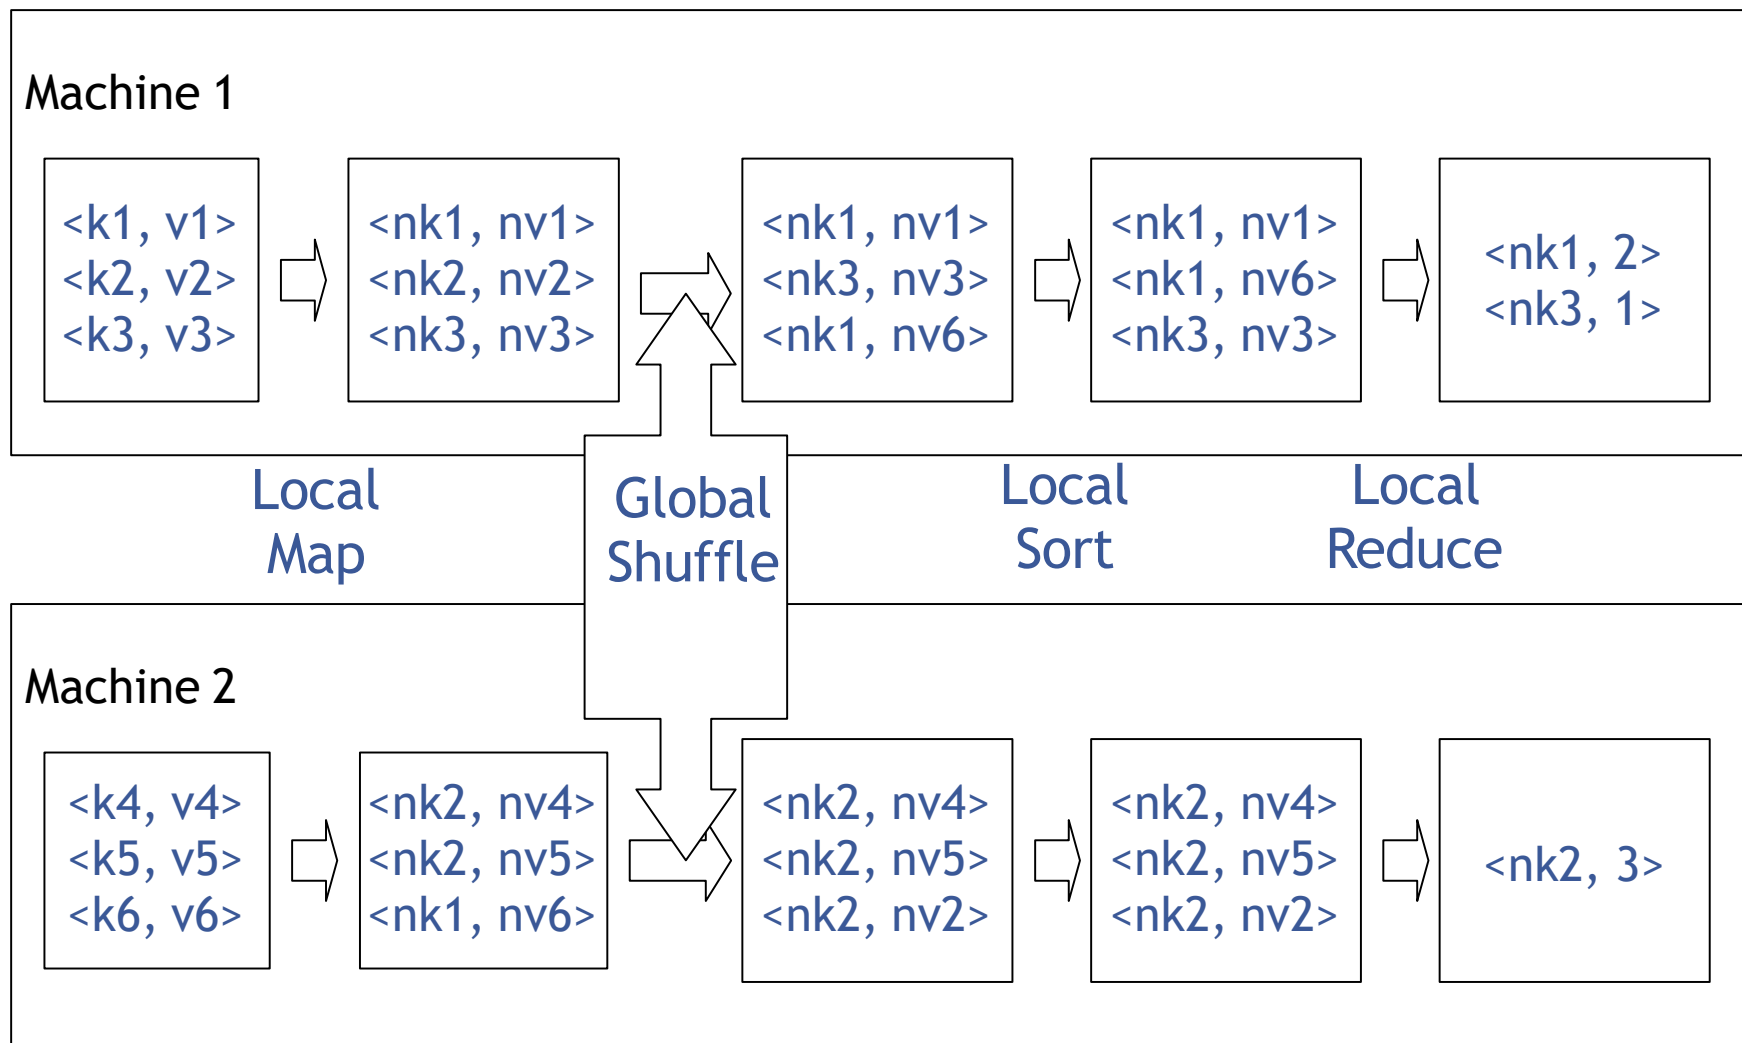
Hive Query Language

- Basic SQL
 - From clause subquery
 - ANSI JOIN (equi-join only)
 - Multi-table Insert
 - Multi group-by
 - Sampling
 - Objects traversal
- Extensibility
 - Pluggable Map-reduce scripts using TRANSFORM

Running Custom Map/Reduce Scripts

```
FROM (  
    FROM pv_users  
    SELECT TRANSFORM(pv_users.userid, pv_users.date) USING  
        'map_script'  
    AS(dt, uid)  
    CLUSTER BY(dt)) map  
INSERT INTO TABLE pv_users_reduced  
    SELECT TRANSFORM(map.dt, map.uid) USING 'reduce_script'  
    AS (date, count);
```

(Simplified) Map Reduce Review



Hive QL - Join

page_view

| pageid | userid | time |
|--------|------------|---------|
| 1 | 111 | 9:08:01 |
| 2 | 111 | 9:08:13 |
| 1 | 222 | 9:08:14 |

X

user

| userid | age | gender |
|------------|-----|--------|
| 111 | 25 | female |
| 222 | 32 | male |

=

pv_users

| pageid | age |
|--------|-----|
| 1 | 25 |
| 2 | 25 |
| 1 | 32 |

- SQL:
INSERT INTO TABLE pv_users
SELECT pv.pageid, u.age
FROM page_view pv JOIN user u ON (pv.userid = u.userid);

Hive QL - Join in Map Reduce

page_view

| page id | user id | time |
|---------|------------|---------|
| 1 | 111 | 9:08:01 |
| 2 | 111 | 9:08:13 |
| 1 | 222 | 9:08:14 |



| key | value |
|-----|----------------|
| 111 | < 1 ,1> |
| 111 | < 1 ,2> |
| 222 | < 1 ,1> |

Map

user

| user id | age | gender |
|------------|-----|--------|
| 111 | 25 | female |
| 222 | 32 | male |



| key | value |
|-----|-----------------|
| 111 | < 2 ,25> |
| 222 | < 2 ,32> |

Shuffle
Sort

| key | value |
|-----|-----------------|
| 111 | < 1 ,1> |
| 111 | < 1 ,2> |
| 111 | < 2 ,25> |



Reduce

| key | value |
|-----|-----------------|
| 222 | < 1 ,1> |
| 222 | < 2 ,32> |



Joins

- Outer Joins

```
INSERT INTO TABLE pv_users
```

```
SELECT pv.*, u.gender, u.age
```

```
FROM page_view pv FULL OUTER JOIN user u ON (pv.userid = u.id)
```

```
WHERE pv.date = 2008-03-03;
```


Join To Map Reduce

- Only Equality Joins with conjunctions supported
- Future
 - Pruning of values send from map to reduce on the basis of projections
 - Make Cartesian product more memory efficient
 - Map side joins
 - Hash Joins if one of the tables is very small
 - Exploit pre-sorted data by doing map-side merge join

Hive Optimizations

- Merge Sequential Map Reduce Jobs

A

| key | a v |
|-----|-----|
| 1 | 111 |

B

| key | b v |
|-----|-----|
| 1 | 222 |

Map Reduce

AB

| key | a v | b v |
|-----|-----|-----|
| 1 | 111 | 222 |

C

| key | c v |
|-----|-----|
| 1 | 333 |

Map Reduce

ABC

| key | a v | b v | c v |
|-----|-----|-----|-----|
| 1 | 111 | 222 | 333 |


■ SQL:

- FROM (a join b on a.key = b.key) join c on a.key = c.key
SELECT ...

Hive QL - Group By

pv_users

| pageid | age |
|--------|-----|
| 1 | 25 |
| 2 | 25 |
| 1 | 32 |
| 2 | 25 |




| pageid | age | count |
|--------|-----|-------|
| 1 | 25 | 1 |
| 2 | 25 | 2 |
| 1 | 32 | 1 |

```
SELECT pageid, age, count(1)
FROM pv_users
GROUP BY pageid, age;
```

Hive QL - Group By in Map Reduce

pv_users


| pageid | age |
|--------|-----|
| 1 | 25 |
| 2 | 25 |



| key | value |
|--------|-------|
| <1,25> | 1 |
| <2,25> | 1 |

Map


| pageid | age |
|--------|-----|
| 1 | 32 |
| 2 | 25 |



| key | value |
|--------|-------|
| <1,32> | 1 |
| <2,25> | 1 |


Shuffle
Sort

| key | value |
|--------|-------|
| <1,25> | 1 |
| <1,32> | 1 |



Reduce

| key | value |
|--------|-------|
| <2,25> | 1 |
| <2,25> | 1 |




| pa |
|----|
| |
| |

| pa |
|----|
| |

Hive QL - Group By with Distinct

page_view

| pageid | userid | time |
|--------|--------|---------|
| 1 | 111 | 9:08:01 |
| 2 | 111 | 9:08:13 |
| 1 | 222 | 9:08:14 |
| 2 | 111 | 9:08:20 |



| pageid | count_distinct_userid |
|--------|-----------------------|
| 1 | 2 |
| 2 | 1 |

```
SELECT pageid, COUNT(DISTINCT userid)
FROM page_view GROUP BY pageid
```

Hive QL - Group By with Distinct in Map Reduce

page_view

| pa ge id | us er id | time |
|----------|----------|---------|
| 1 | 111 | 9:08:01 |
| 2 | 111 | 9:08:13 |

| pa ge id | us er id | time |
|----------|----------|---------|
| 1 | 222 | 9:08:14 |
| 2 | 111 | 9:08:20 |

Shuffle
and
Sort

| ke y | v |
|---------|---|
| <1,111> | |
| <2,111> | |
| <2,111> | |

| ke y | v |
|---------|---|
| <1,222> | |

Reduce

| pa ge id | count |
|----------|-------|
| 1 | 1 |
| 2 | 1 |

| pa ge id | count |
|----------|-------|
| 1 | 1 |

Group by Future optimizations

- Map side partial aggregations
 - Hash Based aggregates
 - Exploit pre-sorted data for distinct counts
- Partial aggregations in Combiner
- Be smarter about how to avoid multiple stage
- Exploit table/column statistics for deciding strategy

Inserts into Files, Tables and Local Files

```
FROM pv_users
INSERT INTO TABLE pv_gender_sum
  SELECT pv_users.gender, count_distinct(pv_users.userid)
  GROUP BY(pv_users.gender)
INSERT INTO DIRECTORY '/user/facebook/tmp/pv_age_sum.dir'
  SELECT pv_users.age, count_distinct(pv_users.userid)
  GROUP BY(pv_users.age)
INSERT INTO LOCAL DIRECTORY '/home/me/pv_age_sum.dir'
  FIELDS TERMINATED BY ',' LINES TERMINATED BY \013
  SELECT pv_users.age, count_distinct(pv_users.userid)
  GROUP BY(pv_users.age);
```


Future Work

- Cost-based optimization
- Multiple interfaces (JDBC...)
- Performance Comparisons with similar work (PIG)
- SQL Compliance (order by, nested queries...)
- Integration with BI tools
- Data Compression
 - Columnar storage schemes
 - Exploit lazy/functional Hive field retrieval interfaces
- Better data locality
 - Co-locate hash partitions on same rack
 - Exploit high intra-rack bandwidth for merge joins

Hive Performance

- full table aggregate (not grouped)
- Input data size: 1,407,867,660 (32 files)
- count in mapper and 2 map-reduce jobs for sum
 - time taken 30 seconds
 - Test cluster: 10 nodes

```
from (  
    from test t select transform (t.userid) as (cnt) using myCount'  
    ) mout  
select sum(mout.cnt);
```

Hadoop Challenges @ Facebook

- QOS/Isolation:
 - Big jobs can hog the cluster
 - JobTracker memory as limited resource
 - Limit memory impact of runaway tasks
 - ➔ Fair Scheduler (Matei)
- Protection
 - What if a software bug corrupts the NameNode transaction log/image?
 - ➔ HDFS SnapShots (Dhruba)
- Data Archival
 - Not all data is hot and needs colocation with Compute
 - ➔ HDFS Symlinks (Dhruba) Data Archival
- Performance
 - Really hard to understand what bottlenecks are

Conclusion

- Available as a contrib project in hadoop
 - <http://svn.apache.org/repos/asf/hadoop/core/>
 - Checkout src/contrib/hive from trunk (works against 0.19 onwards)
- Latest distributions (including for hadoop-0.17) at:
<http://mirror.facebook.com/facebook/hive/>
- People:
 - Suresh Anthony
 - Zheng Shao
 - Prasad Chakka
 - Pete Wyckoff
 - Namit Jain
 - Raghu Murthy
 - Joydeep Sen Sarma
 - Ashish Thusoo